Assignment -2
Data Visualization and Preprocessing

Assignment Date	26 September 2022
Team ID	PNT2022TMID05805
Project Name	AI BASED DISCOURSE FOR BANKING INDUSTRY
Student Name	ELANGO N S
Student Roll Number	1905022
Maximum Marks	2 Marks

Question-1. Download dataset

Solution:

RowNumb	Customer	Surname	CreditSco	Geograph	Gender	Age	Tenure	Balance	NumOfPr	HasCrCard	IsActiveM	Estimated	Exited
1	15634602	Hargrave	619	France	Female	42	2	0	1	1	1	101348.9	1
2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.6	0
3	15619304	Onio	502	France	Female	42	8	159660.8	3	1	0	113931.6	1
4	15701354	Boni	699	France	Female	39	1	0	2	0	0	93826.63	0
5	15737888	Mitchell	850	Spain	Female	43	2	125510.8	1	1	1	79084.1	0
6	15574012	Chu	645	Spain	Male	44	8	113755.8	2	1	0	149756.7	1
7	15592531	Bartlett	822	France	Male	50	7	0	2	1	1	10062.8	0
8	15656148	Obinna	376	Germany	Female	29	4	115046.7	4	1	0	119346.9	1
9	15792365	He	501	France	Male	44	4	142051.1	2	0	1	74940.5	0
10	15592389	H?	684	France	Male	27	2	134603.9	1	1	1	71725.73	0
11	15767821	Bearce	528	France	Male	31	6	102016.7	2	0	0	80181.12	0
12	15737173	Andrews	497	Spain	Male	24	3	0	2	1	0	76390.01	0
13	15632264	Kay	476	France	Female	34	10	0	2	1	0	26260.98	0
14	15691483	Chin	549	France	Female	25	5	0	2	0	0	190857.8	0
15	15600882	Scott	635	Spain	Female	35	7	0	2	1	1	65951.65	0
16	15643966	Goforth	616	Germany	Male	45	3	143129.4	2	0	1	64327.26	0
17	15737452	Romeo	653	Germany	Male	58	1	132602.9	1	1	0	5097.67	1
18	15788218	Henderso	549	Spain	Female	24	9	0	2	1	1	14406.41	0
19	15661507	Muldrow	587	Spain	Male	45	6	0	1	0	0	158684.8	0
20	15568982	Hao	726	France	Female	24	6	0	2	1	1	54724.03	0
21	15577657	McDonald	732	France	Male	41	8	0	2	1	1	170886.2	0
22	15597945	Dellucci	636	Spain	Female	32	8	0	2	1	0	138555.5	0
23	15699309	Gerasimo	510	Spain	Female	38	4	0	1	1	0	118913.5	1
24	15725737	Mosman	669	France	Male	46	3	0	2	0	1	8487.75	0
25	15625047	Yen	846	France	Female	38	5	0	1	1	1	187616.2	0
26	15738191	Maclean	577	France	Male	25	3	0	2	0	1	124508.3	0
27	15736816	Young	756	Germany	Male	36	2	136815.6	1	1	1	170042	0
28	15700772	Nebechi	571	France	Male	44	9	0	2	0	0	38433.35	0
29	15728693	McWillian	574	Germany	Female	43	3	141349.4	1	1	1	100187.4	0
30	15656300	Lucciano	411	France	Male	29	0	59697.17	2	1	1	53483.21	0
31	15589475	Azikiwe	591	Spain	Female	39	3	0	3	1	0	140469.4	1
32	15706552	Odinakach	533	France	Male	36	7	85311.7	1	0	1	156731.9	0
33	15750181	Sanderson	553	Germany	Male	41	9	110112.5	2	0	0	81898.81	0
34	15659428	Maggard	520	Spain	Female	42	6	0	2	1	1	34410.55	0
35	15732963	Clements	722	Spain	Female	29	9	0	2	1	1	142033.1	0
36	15794171	Lombardo	475	France	Female	45	0	134264	1	1	0	27822.99	1
37	15788448	Watson	490	Spain	Male	31	3	145260.2	1	0	1	114066.8	0
38	15729599	Lorenzo	804	Spain	Male	33	7	76548.6	1	0	1	98453.45	0
39	15717426	Armstrong	850	France	Male	36	7	0	1	1	1	40812.9	0
40	15585768	Cameron	582	Germany	Male	41	6	70349.48	2	0	1	178074	0

Question-2. Load the dataset

import numpy as np import
pandas as pd import seaborn
as sns import
matplotlib.pyplot as plt import
sklearn
data = pd.read_csv(r'Churn_Modelling.csv') df.head

```
<bound method NDFrame.head of        RowNumber  CustomerId    Surname  CreditScore Geography  Gender  Age
0              1    15634602   Hargrave          619    France  Female   42
1              2    15647311       Hill          608     Spain  Female   41
2              3    15619304       Onio          502    France  Female   42
3              4    15701354       Boni          699    France  Female   39
4              5    15737888   Mitchell          850     Spain  Female   43
...          ...         ...        ...          ...       ...     ...  ...
9995        9996    15606229   Obijiaku          771    France    Male   39
9996        9997    15569892  Johnstone          516    France    Male   35
9997        9998    15584532        Liu          709    France  Female   36
9998        9999    15682355   Sabbatini          772   Germany    Male   42
9999       10000    15628319     Walker          792    France  Female   28

      Tenure    Balance  NumOfProducts  HasCrCard  IsActiveMember  \
0          2       0.00              1          1               1
1          1   83807.86              1          0               1
2          8  159660.80              3          1               0
3          1       0.00              2          0               0
4          2  125510.82              1          1               1
...      ...        ...            ...        ...             ...
9995       5       0.00              2          1               0
9996      10   57369.61              1          1               1
9997       7       0.00              1          0               1
9998       3   75075.31              2          1               0
9999       4  130142.79              1          1               0

      EstimatedSalary  Exited
0           101348.88       1
1           112542.58       0
2           113931.57       1
3            93826.63       0
4            79084.10       0
...               ...     ...
9995         96270.64       0
9996        101699.77       0
9997         42085.58       1
9998         92888.52       1
9999         38190.78       0

[10000 rows x 14 columns]>
```

Question-3. Perform Below Visualizations.

3.1 Univariate Analysis

Solution:

sns.displot(data['CreditScore'])


```
<seaborn.axisgrid.FacetGrid at 0x7fbdf8e05510>
```

sns.histplot(data['CreditScore'])


```
<matplotlib.axes._subplots.AxesSubplot at 0x7fbdf8c6b8d0>
```

sns.boxplot(x = data['CreditScore'])

<matplotlib.axes._subplots.AxesSubplot at 0x7fbdf8bc6910>


sns.distplot(data['Age'])

<matplotlib.axes._subplots.AxesSubplot at 0x7fbe0d180550>

sns.histplot(data['Age'])

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fbe0d15f110>
```

sns.boxplot(data['Age'])


```
<matplotlib.axes._subplots.AxesSubplot at 0x7fbe0ca3c0d0>
```

3.2 Bivariate Analysis

Solution:

plt.figure(figsize=(7,7))
sns.lineplot(data = data, x = 'Tenure', y = 'CreditScore')

`<matplotlib.axes._subplots.AxesSubplot at 0x7fbdf8aaa350>`

plt.figure(figsize=(10,10))
sns.barplot(data = data, x = 'CreditScore', y ='EstimatedSalary')

plt.figure(figsize=(10,10))
sns.barplot(data = data, x = 'CreditScore', y = 'Tenure')


plt.figure(figsize=(10,10))
sns.lineplot(data['Age'], data['EstimatedSalary'])

plt.figure(figsize=(17,17))
**sns.barplot(data['Age'], data['EstimatedSalary'])**

**sns.scatterplot(data = data, x = 'CreditScore', y = 'Age')**


3.3 Multivariate Analysis

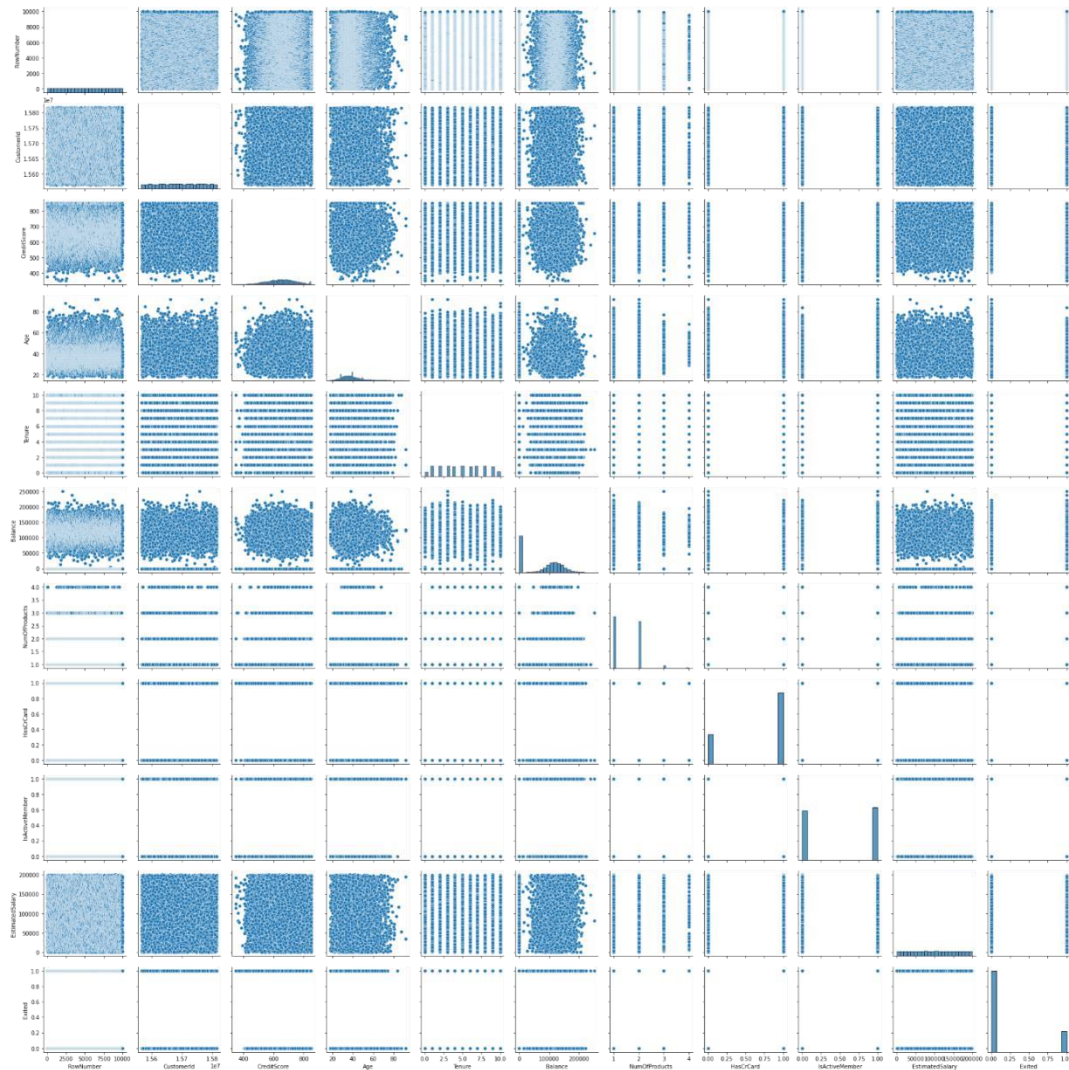**Solution:** sns.scatterplot(data = data, x = 'CreditScore', y = 'Balance', hue = 'Gender')

sns.scatterplot(data['Tenure'], data['CreditScore'], hue = data['Gender'])

sns.scatterplot(data['Age'], data['Balance'], hue = data['Gender'])

sns.pairplot(data)

Question-4. Perform descriptive statistics on the dataset.

Solution: data.mean(numeric_only

= True)

```
RowNumber          5.000500e+03
CustomerId         1.569094e+07
CreditScore        6.505288e+02
Age                3.892180e+01
Tenure             5.012800e+00
Balance            7.648589e+04
NumOfProducts      1.530200e+00
HasCrCard          7.055000e-01
IsActiveMember     5.151000e-01
EstimatedSalary    1.000902e+05
Exited             2.037000e-01
dtype: float64
```

data.median(numeric_only = True)

```
RowNumber          5.000500e+03
CustomerId         1.569074e+07
CreditScore        6.520000e+02
Age                3.700000e+01
Tenure             5.000000e+00
Balance            9.719854e+04
NumOfProducts      1.000000e+00
HasCrCard          1.000000e+00
IsActiveMember     1.000000e+00
EstimatedSalary    1.001939e+05
Exited             0.000000e+00
dtype: float64
```

data['CreditScore'].mode()

```
0    850
dtype: int64
```

data['EstimatedSalary'].mode()

```
0    24924.92
dtype: float64
```

data['HasCrCard'].unique()

```
array([1, 0])
```

data['Tenure'].unique()

```
array([ 2,  1,  8,  7,  4,  6,  3, 10,  5,  9,  0])
```

data.std(numeric_only=True)

```
RowNumber          2886.895680
CustomerId        71936.186123
CreditScore          96.653299
Age                  10.487806
Tenure                2.892174
Balance           62397.405202
NumOfProducts         0.581654
HasCrCard             0.455840
IsActiveMember        0.499797
EstimatedSalary   57510.492818
Exited                0.402769
dtype: float64
```

data.describe()

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.00000	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.239881	0.203700
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.499797	57510.492818	0.402769
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000	11.580000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000000	51002.110000	0.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.000000	100193.915000	0.000000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000	1.000000	149388.247500	0.000000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.00000	1.000000	199992.480000	1.000000

data['Tenure'].value_counts()

```
 2     1048
 1     1035
 7     1028
 8     1025
 5     1012
 3     1009
 4      989
 9      984
 6      967
 10     490
 0      413
Name: Tenure, dtype: int64
```

Question-5. Handle the Missing values.

Solution: data.isnull().any()

```
RowNumber          False
CustomerId         False
Surname            False
CreditScore        False
Geography          False
Gender             False
Age                False
Tenure             False
Balance            False
NumOfProducts      False
HasCrCard          False
IsActiveMember     False
EstimatedSalary    False
Exited             False
dtype: bool
```

data.isnull().sum()

```
RowNumber          0
CustomerId         0
Surname            0
CreditScore        0
Geography          0
Gender             0
Age                0
Tenure             0
Balance            0
NumOfProducts      0
HasCrCard          0
IsActiveMember     0
EstimatedSalary    0
Exited             0
dtype: int64
```
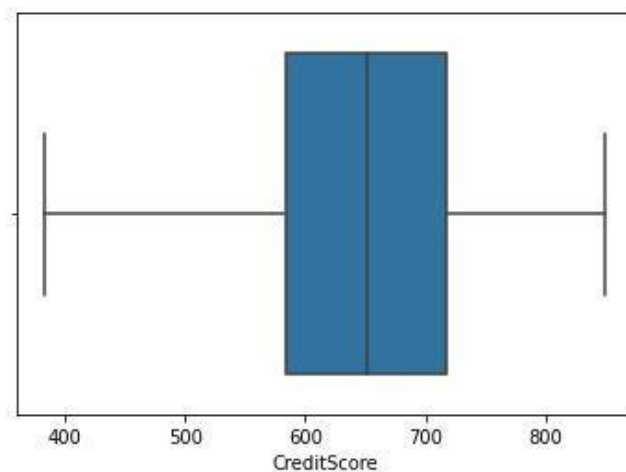
Question-6. Find the outliers and replace the outliers

Solution: sns.boxplot(data['CreditScore'])   #Outlier detection

- box plot

fig, ax = plt.subplots(figsize = (5,3))   #Outlier detection - Scatter plot
ax.scatter(data['Balance'], data['Exited'])

 # x-axis label
ax.set_xlabel('Balance')

 # y-axis label ax.set_ylabel('Exited')
plt.show()

sns.boxplot(x=data['Balance'])

**from scipy import stats   #Outlier detection – zscore zscore
= np.abs(stats.zscore(data['CreditScore'])) print(zscore)
print('No. of Outliers : ', np.shape(np.where(zscore>3)))**

```
0        0.332952
1        0.447540
2        1.551761
3        0.500422
4        2.073415
          ...
9995     1.250458
9996     1.405920
9997     0.604594
9998     1.260876
9999     1.469219
Name: CreditScore, Length: 10000, dtype: float64
No. of Outliers :  (1, 0)
```

q = data.quantile([0.75,0.25]) q

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0.75	7500.25	15753233.75	2238.25	718.0	1.0	1.0	44.0	7.0	127644.24	2.0	1.0	1.0	149388.2475	0.0
0.25	2500.75	15628528.25	773.75	584.0	0.0	0.0	32.0	3.0	0.00	1.0	0.0	0.0	51002.1100	0.0

iqr = q.iloc[0] - q.iloc[1] iqr

```
RowNumber          4999.5000
CustomerId       124705.5000
Surname            1464.5000
CreditScore         134.0000
Geography             1.0000
Gender                1.0000
Age                  12.0000
Tenure                4.0000
Balance          127644.2400
NumOfProducts         1.0000
HasCrCard             1.0000
IsActiveMember        1.0000
EstimatedSalary   98386.1375
Exited                0.0000
dtype: float64
```

u = q.iloc[0] + (1.5*iqr)   u

```
RowNumber        1.499950e+04
CustomerId       1.594029e+07
Surname          4.435000e+03
CreditScore      9.190000e+02
Geography        2.500000e+00
Gender           2.500000e+00
Age              6.200000e+01
Tenure           1.300000e+01
Balance          3.191106e+05
NumOfProducts    3.500000e+00
HasCrCard        2.500000e+00
IsActiveMember   2.500000e+00
EstimatedSalary  2.969675e+05
Exited           0.000000e+00
dtype: float64
```

l = q.iloc[1] - (1.5*iqr)
l

```
RowNumber        -4.998500e+03
CustomerId        1.544147e+07
Surname          -1.423000e+03
CreditScore       3.830000e+02
Geography        -1.500000e+00
Gender           -1.500000e+00
Age               1.400000e+01
Tenure           -3.000000e+00
Balance          -1.914664e+05
NumOfProducts    -5.000000e-01
HasCrCard        -1.500000e+00
IsActiveMember   -1.500000e+00
EstimatedSalary  -9.657710e+04
Exited            0.000000e+00
dtype: float64
```

Q1 = data['EstimatedSalary'].quantile(0.25)   #Outlier detection - IQR

```python
Q3 = data['EstimatedSalary'].quantile(0.75) iqr = Q3 - Q1 print(iqr)
upper=Q3 + 1.5 * iqr lower=Q1 - 1.5 * iqr count =
np.size(np.where(data['EstimatedSalary'] >upper)) count = count
+ np.size(np.where(data['EstimatedSalary'] <lower)) print('No. of
outliers : ', count)
```

```
98386.1375
No. of outliers :  0
```

```python
data['CreditScore'] = np.where(np.logical_or(data['CreditScore']>900, data['CreditScore']<383), 65
0, data['CreditScore']) sns.boxplot(data['CreditScore'])
```



```python
upper = data.Age.mean() + (3 * data.Age.std())    #Outlier detection - 3 sigma
lower = data.Age.mean() - (3 * data.Age.std()) columns = data[ ( data['Age']
> upper ) | ( data['Age']<lower ) ] print('Upper range : ', upper) print('Lower
range : ', lower) print('No. of Outliers : ', len(columns))
```

```
Upper range :  70.38521935511383
Lower range :  7.458380644886169
No. of Outliers :  133
```

```python
columns = ['EstimatedSalary', 'Age', 'Balance', 'NumOfProducts', 'Tenure', 'CreditScore']
#After outlier removal
```

```python
for i in columns:
  Q1 = data[i].quantile(0.25)
Q3 = data[i].quantile(0.75)
iqr = Q3 - Q1   upper=Q3 +
1.5 * iqr   lower=Q1 - 1.5 *
iqr
```

```
  count = np.size(np.where(data[i] >upper))   count
= count + np.size(np.where(data[i] <lower))
print('No. of outliers in ', i, ' : ', count)
```

```
No. of outliers in  EstimatedSalary  :  0
No. of outliers in  Age  :  0
No. of outliers in  Balance  :  0
No. of outliers in  NumOfProducts  :  0
No. of outliers in  Tenure  :  0
No. of outliers in  CreditScore  :  0
```

Question-7. Check for Categorical columns and perform encoding

Solution:

```
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
le = LabelEncoder() oneh = OneHotEncoder()
data['Surname'] = le.fit_transform(data['Surname']) data['Gender']
= le.fit_transform(data['Gender']) data['Geography'] =
le.fit_transform(data['Geography']) data.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	1115	619	0	0	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	1177	608	2	0	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	2040	502	0	0	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	289	699	0	0	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	1822	850	2	0	43	2	125510.82	1	1	1	79084.10	0

Question-8. Split the data into dependent and independent variables split the data in X and Y

Solution:

```
x        # independent values (
inputs)    x = data.iloc[:, 0:13]
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	1115	619	0	0	42	2	0.00	1	1	1	101348.88
1	2	15647311	1177	608	2	0	41	1	83807.86	1	0	1	112542.58
2	3	15619304	2040	502	0	0	42	8	159660.80	3	1	0	113931.57
3	4	15701354	289	699	0	0	39	1	0.00	2	0	0	93826.63
4	5	15737888	1822	850	2	0	43	2	125510.82	1	1	1	79084.10
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	1999	771	0	1	39	5	0.00	2	1	0	96270.64
9996	9997	15569892	1336	516	0	1	35	10	57369.61	1	1	1	101699.77
9997	9998	15584532	1570	709	0	0	36	7	0.00	1	0	1	42085.58
9998	9999	15682355	2345	772	1	1	42	3	75075.31	2	1	0	92888.52
9999	10000	15628319	2751	792	0	0	28	4	130142.79	1	1	0	38190.78

10000 rows × 13 columns

y   # dependent values
(output)   y = data['Exited']

```
0        1
1        0
2        1
3        0
4        0
        ..
9995     0
9996     0
9997     1
9998     1
9999     0
Name: Exited, Length: 10000, dtype: int64
```

Question-9. Scale the independent variables

Solution:

from sklearn.preprocessing import StandardScaler, MinMaxScaler
sc = StandardScaler() x_scaled = sc.fit_transform(x) x_scaled

```
array([[-1.73187761, -0.78321342, -0.46418322, ...,  0.64609167,
         0.97024255,  0.02188649],
        [-1.7315312 , -0.60653412, -0.3909112 , ..., -1.54776799,
         0.97024255,  0.21653375],
        [-1.73118479, -0.99588476,  0.62898807, ...,  0.64609167,
        -1.03067011,  0.2406869 ],
        ...,
        [ 1.73118479, -1.47928179,  0.07353887, ..., -1.54776799,
         0.97024255, -1.00864308],
        [ 1.7315312 , -0.11935577,  0.98943914, ...,  0.64609167,
        -1.03067011, -0.12523071],
        [ 1.73187761, -0.87055909,  1.4692527 , ...,  0.64609167,
        -1.03067011, -1.07636976]])
```

Question-10. Split x and y into Training and Testing

Solution:

from sklearn.model_selection import train_test_split x_train, x_test, y_train, y_test = train_test_split(x_scaled, y, test_size = 0.3, random_state = 0)

x_train

```
array([[ 0.92889885, -0.79703192, -1.47580983, ...,  0.64609167,
         0.97024255, -0.77021814],
        [ 1.39655257,  0.71431365, -1.58808148, ...,  0.64609167,
        -1.03067011, -1.39576675],
        [-0.4532777 ,  0.96344969, -0.24082173, ..., -1.54776799,
         0.97024255, -1.49965629],
        ...,
        [-0.60119484, -1.62052514, -0.36136603, ...,  0.64609167,
        -1.03067011,  1.41441489],
        [ 1.67853045, -0.37403866,  0.72589622, ...,  0.64609167,
         0.97024255,  0.84614739],
        [-0.78548505, -1.36411841,  1.3829808 , ...,  0.64609167,
        -1.03067011,  0.32630495]])
```

x_train.shape

```
(7000, 13)
```

x_test

```
array([[ 1.52229946, -1.04525042,  1.39834429, ...,  0.64609167,
         0.97024255,  1.61304597],
       [-1.42080128, -0.50381294, -0.78208925, ...,  0.64609167,
        -1.03067011,  0.49753166],
       [-0.90118604, -0.7932923 ,  0.41271742, ...,  0.64609167,
         0.97024255, -0.4235611 ],
       ...,
       [ 1.49216178, -0.14646448,  0.6868966 , ...,  0.64609167,
         0.97024255,  1.17045451],
       [ 1.1758893 , -1.29228727, -1.38481071, ...,  0.64609167,
         0.97024255, -0.50846777],
       [ 0.08088677, -1.38538833,  1.11707427, ...,  0.64609167,
         0.97024255, -1.15342685]])
```

x_test.shape

```
(3000, 13)
```

y_train

```
7681    1
9031    0
3691    0
202     1
5625    0
        ..
9225    0
4859    0
3264    0
9845    0
2732    1
Name: Exited, Length: 7000, dtype: int64
```

y_test

```
9394    0
898     1
2398    0
5906    0
2343    0
        ..
4004    0
7375    0
9307    0
8394    0
5233    1
Name: Exited, Length: 3000, dtype: int64
```